

# Introduction to AWS CloudFormation

---

AWS CloudFormation allows you to manage your AWS infrastructure as code. Instead of manually creating and configuring resources, you can define your infrastructure in a CloudFormation template, which is a text file that describes the desired state of your AWS environment.

## Key Features of CloudFormation

### 1. Infrastructure as Code:

- Define your AWS resources and their configurations in code.
- Example: You can specify security groups, EC2 instances, Elastic IPs, S3 buckets, and load balancers in a CloudFormation template.
- CloudFormation automates the creation and configuration of these resources in the correct order.

### 2. Declarative Approach:

- Specify what you want (e.g., two EC2 instances, a load balancer), and CloudFormation handles the "how" for you.
- The template is a declarative document where you state what your infrastructure should look like.

### 3. Version Control and Reusability:

- CloudFormation templates can be version-controlled using systems like Git.
- Changes to infrastructure are tracked and can be reviewed.

### 4. Cost Management:

- Resources created by CloudFormation are tagged for easy cost tracking.
- Estimate costs from your templates before deployment.

### 5. Automated Environment Management:

- Automate the creation and deletion of environments.
- Ideal for development environments that can be deleted at the end of the day and recreated the next morning, leveraging the pay-as-you-go model of AWS.

### 6. Automated Diagrams:

- Generate diagrams for your infrastructure directly from the templates, useful for architecture documentation.

### 7. Separation of Concerns:

- Use multiple CloudFormation stacks to manage different parts of your infrastructure (e.g., networking, application layers).

## 8. Leveraging Existing Templates:

- Use and adapt pre-existing templates available online to save time and effort.

## How CloudFormation Works

### 1. Template Creation:

- Write CloudFormation templates using either YAML or JSON.
- Define resources, configurations, and dependencies.

### 2. Template Upload:

- Upload your templates to Amazon S3.
- Reference these templates in CloudFormation to create stacks.

### 3. Stack Management:

- A stack represents a collection of AWS resources managed together.
- Update stacks by uploading new versions of templates.
- Deleting a stack removes all resources created by that stack.

## Deploying CloudFormation Templates

### 1. Manual Deployment:

- Create templates using tools like AWS Application Composer or any text editor.
- Deploy templates via the AWS Management Console.

### 2. Automated Deployment:

- Edit templates in a code editor and use the AWS CLI or CI/CD tools for deployment.
- Automate the entire infrastructure deployment process, enhancing productivity and consistency.

## Components of a CloudFormation Template

### 1. Template Format Version:

- Defines the version of the template format (for AWS internal use).

### 2. Description:

- Optional comments about the template.

### 3. Resources:

- The only mandatory section.
- Specifies the AWS resources to be created.

### 4. Parameters:

- Allow for dynamic inputs when creating the stack.

## 5. **Mappings:**

- Define static variables that can be used in the template.

## 6. **Outputs:**

- Provide information about the resources created.

## 7. **Conditionals:**

- Define conditions for resource creation.

## 8. **References and Functions:**

- Help with the template logic and resource dependencies.

By using AWS CloudFormation, you can efficiently manage your infrastructure, ensure consistency across environments, and leverage the full power of AWS automation and scalability.

# Practicing with AWS CloudFormation

---

To start using AWS CloudFormation, follow these steps:

## 1. Select the Region:

- Ensure you are in the US East (N. Virginia) region, `us-east-1`, as templates are designed for this region, especially when using AMI IDs.

## 2. Access CloudFormation:

- Open the CloudFormation service. If you have existing stacks, they will be listed here; otherwise, you will see zero stacks.

## 3. Create a Stack:

- Click on "Create stack". You have several options to create a stack:
  - Use an existing template.
  - Use a sample template.
  - Build directly from Application Composer.

## 4. Use a Sample Template:

- Choose a sample template like "Multi\_AZ\_Simple" or "WordPress blog".
- Instead of launching it, click to view it in Application Composer.

## 5. View in Application Composer:

- Application Composer provides a visual and code representation of the template.
- Click on "Template" to switch between YAML and JSON formats. YAML is preferred for readability.

## 6. Explore the Template:

- The template file outlines the resources and their configurations.
- Visual representation in Canvas shows components such as `WebServerSecurityGroup`, `LaunchConfig`, `WebServerGroup`, and database instances.
- Clicking on these components shows their configurations as defined in the CloudFormation template.

## Deploying a Custom Template

### 1. Open the Custom YAML File:

- Locate and open `0-just-ec2.yaml` in your code editor.
- This file contains the definition for creating a single EC2 instance.

### 2. Understanding the YAML File:

- The file defines a mandatory `Resources` section.
- It includes one resource named `MyInstance` of type `AWS::EC2::Instance`.

- Properties specified include:
  - `AvailabilityZone`: `us-east-1a`
  - `ImageId`: a specific AMI ID
  - `InstanceType`: `t2.micro`

### 3. Upload and Create the Stack:

- In the CloudFormation console, choose to upload a template file and select `0-just-ec2.yaml`.
- Name your stack `EC2InstanceDemo`.
- Skip additional settings like tags and permissions for now and proceed to create the stack.
- CloudFormation uploads the template to S3 and references it for stack creation.

### 4. Monitor the Stack Creation:

- After submitting, monitor the creation progress in the "Events" tab.
- You will see events indicating the creation status, such as `CREATE_IN_PROGRESS` and `CREATE_COMPLETE`.

### 5. Verify the Created Resource:

- Once complete, go to the "Resources" tab and click on the Physical ID of `MyInstance`.
- This directs you to the EC2 console to view the instance.
- Verify the instance details:
  - Instance type: `t2.micro`
  - Availability zone: `us-east-1a`
  - AMI ID: matches the specified ID
- Check the Tags tab to see tags added by CloudFormation, including stack ID, logical ID, and stack name.

## CloudFormation Stack Information

- **Stack Info:** General information about the stack.
- **Events:** Detailed log of events during the stack creation process.
- **Resources:** List of resources created by the stack.
- **Outputs:** Any defined outputs from the stack.
- **Parameters:** Parameters passed to the template (none in this case).
- **Template:** The original template file content.

This hands-on exercise demonstrates the power of CloudFormation in automating the creation and management of AWS resources using declarative templates.

# Example of sample of a cloudFormation Template

---

Here's a simple example of a CloudFormation template in YAML that creates an EC2 instance and an S3 bucket. This template includes basic configurations and demonstrates the use of parameters and outputs.

```
AWSTemplateFormatVersion: '2010-09-09'
Description: Simple CloudFormation template to create an EC2 instance and an S3 bucket.

Parameters:
  InstanceType:
    Description: EC2 instance type
    Type: String
    Default: t2.micro
    AllowedValues: [t2.micro, t2.small, t2.medium]
    ConstraintDescription: Must be a valid EC2 instance type.

Resources:
  MyEC2Instance:
    Type: 'AWS::EC2::Instance'
    Properties:
      InstanceType: !Ref InstanceType
      ImageId: ami-0c55b159cbfafa1f0 # This is a placeholder; use a valid AMI
      ID for your region
      KeyName: !Ref KeyName
      SecurityGroups:
        - !Ref InstanceSecurityGroup
      Tags:
        - Key: Name
          Value: MyEC2Instance

  InstanceSecurityGroup:
    Type: 'AWS::EC2::SecurityGroup'
    Properties:
      GroupDescription: Enable SSH access via port 22
      SecurityGroupIngress:
        - IpProtocol: tcp
          FromPort: '22'
          ToPort: '22'
          CidrIp: 0.0.0.0/0

  MyS3Bucket:
    Type: 'AWS::S3::Bucket'
    Properties:
      BucketName: !Sub '${AWS::StackName}-bucket'

Outputs:
  InstanceId:
    Description: ID of the newly created EC2 instance
    Value: !Ref MyEC2Instance
```

```
BucketName:
  Description: Name of the newly created S3 bucket
  Value: !Ref MyS3Bucket
```

## Explanation

- **AWSTemplateFormatVersion:** Specifies the version of the AWS CloudFormation template format.
- **Description:** A brief description of what the template does.
- **Parameters:**
  - **InstanceType:** Parameter to specify the type of EC2 instance. It has a default value and a set of allowed values.
- **Resources:**
  - **MyEC2Instance:** Defines an EC2 instance with properties like instance type, AMI ID, key name, security groups, and tags.
  - **InstanceSecurityGroup:** Defines a security group allowing SSH access (port 22) from anywhere.
  - **MyS3Bucket:** Defines an S3 bucket with a name based on the stack name.
- **Outputs:**
  - **InstanceId:** Outputs the ID of the created EC2 instance.
  - **BucketName:** Outputs the name of the created S3 bucket.

## Usage

1. Save the above YAML content to a file, for example, `simple-template.yaml`.
2. Go to the AWS CloudFormation console and create a new stack by uploading this template file.
3. Follow the prompts to provide any required parameters and create the stack.
4. Once the stack creation is complete, you can check the outputs to see the EC2 instance ID and the S3 bucket name.

# Updating a CloudFormation Stack

---

To update a CloudFormation stack, follow these steps:

## 1. Select the Current Stack:

- Go to the AWS CloudFormation console and select the stack you want to update.
- Click on "Update" to start the update process.

## 2. Replace the Current Template:

- You can either use the current template or replace it with a new one. For this update, we'll replace the current template.
- Go to your code editor and open the new template file, for example, `one-ec2-with-sg-eip.yaml`.

## 3. New Template Overview:

- The new template includes additional parameters and resources. For example, it might have parameters for security group descriptions, an elastic IP resource, and multiple security groups.
- Example new template structure:

```
AWSTemplateFormatVersion: '2010-09-09'
Description: Updated CloudFormation template to create an EC2
instance, security groups, and an Elastic IP.

Parameters:
  SecurityGroupDescription:
    Description: Description for the security group
    Type: String

Resources:
  MyElasticIP:
    Type: 'AWS::EC2::EIP'

  MyInstance:
    Type: 'AWS::EC2::Instance'
    Properties:
      InstanceType: t2.micro
      ImageId: ami-0c55b159cbfafa1f0 # Use a valid AMI ID
      SecurityGroups:
        - !Ref InstanceSecurityGroup
        - !Ref SSHSecurityGroup
      Tags:
        - Key: Name
          Value: MyUpdatedEC2Instance

  SSHSecurityGroup:
    Type: 'AWS::EC2::SecurityGroup'
```



```

Properties:
  GroupDescription: Allow SSH access
  SecurityGroupIngress:
    - IpProtocol: tcp
      FromPort: '22'
      ToPort: '22'
      CidrIp: 0.0.0.0/0

InstanceSecurityGroup:
  Type: 'AWS::EC2::SecurityGroup'
  Properties:
    GroupDescription: !Ref SecurityGroupDescription
    SecurityGroupIngress:
      - IpProtocol: tcp
        FromPort: '22'
        ToPort: '22'
        CidrIp: 0.0.0.0/0
      - IpProtocol: tcp
        FromPort: '80'
        ToPort: '80'
        CidrIp: 0.0.0.0/0

```

#### 4. Upload the New Template:

- Upload the new template file ([one-ec2-with-sg-eip.yaml](#)).
- Enter the required parameters, for example, the security group description: "This is a cool security group".

#### 5. Review Changes:

- Review the change set preview to see what changes will be made:
  - Adding an Elastic IP
  - Adding an SSH security group
  - Adding a server security group
  - Replacing the existing EC2 instance

#### 6. Submit the Update:

- Submit the update. CloudFormation will begin the update process, creating new resources and replacing the EC2 instance if necessary.
- Monitor the events to see the progress of the update.

#### 7. Verify the Update:

- Once the update is complete, check the resources:
  - Verify that the new EC2 instance is running.
  - Check the Elastic IP and confirm it is associated with the new EC2 instance.
  - Verify the security groups and their inbound rules.
  - Confirm that the security group descriptions match the parameters provided.

## 8. Clean Up Resources:

- If you want to delete the stack and clean up resources, go to the CloudFormation console and delete the stack.
- CloudFormation will handle the deletion of all resources in the correct order.

## Summary

By following these steps, you can update a CloudFormation stack by replacing the current template with a new one. This process involves uploading the new template, reviewing the changes, and monitoring the update process. CloudFormation ensures that resources are created, updated, and deleted in the correct order, making the management of AWS resources efficient and reliable.

# Introduction to YAML for CloudFormation

---

YAML (YAML Ain't Markup Language) is a human-readable data serialization standard that is widely used for configuration files and in applications where data is being stored or transmitted. YAML is particularly popular in the context of CloudFormation for defining AWS resources due to its readability and ease of use compared to JSON.

## Basic Structure of YAML

YAML documents are composed of key-value pairs, where each key is associated with a value. The values can be simple data types or more complex structures, such as nested objects and arrays.

### Key-Value Pairs

- **Basic Key-Value Pair:**

```
invoice: 34843
date: 2024-06-15
```

- `invoice` is a key with a value of `34843` (a number).
- `date` is a key with a value of `2024-06-15` (a string).

### Nested Objects

- **Nested Objects:**

```
bill-to:
  given: Chris
  family: Dumars
  address:
    lines: |
      458 Walkman Dr.
      Suite #292
    city: Royal Oak
    state: MI
    postal: 48046
```

- `bill-to` contains a nested object with keys `given`, `family`, and `address`.
- `address` itself is another nested object containing `lines`, `city`, `state`, and `postal`.

### Arrays

- **Arrays:**

```
products:
  - SKU: BL394D
    quantity: 4
    description: Basketball
    price: 450.00
  - SKU: BL4438H
    quantity: 1
    description: Super Hoop
    price: 2392.00
```

- `products` is an array of objects, each containing `SKU`, `quantity`, `description`, and `price`.

## Multi-line Strings

- **Multi-line Strings:**

```
address:
  lines: |
    458 Walkman Dr.
    Suite #292
```

- The `|` symbol indicates that the value of `lines` spans multiple lines.

## Comments

- **Comments:**

```
# This is a comment
invoice: 34843 # Inline comment
```

- Comments in YAML start with `#` and can be placed on their own line or inline.

## Example CloudFormation Template in YAML

Below is an example of a CloudFormation template written in YAML:

```
AWSTemplateFormatVersion: '2010-09-09'
Description: A simple EC2 instance

Parameters:
  InstanceType:
    Description: Type of EC2 instance
    Type: String
    Default: t2.micro
    AllowedValues:
```

```

- t2.micro
- t2.small
- t2.medium

Resources:
  MyInstance:
    Type: 'AWS::EC2::Instance'
    Properties:
      InstanceType: !Ref InstanceType
      ImageId: ami-0c55b159cbfafa1f0 # Example AMI ID
      SecurityGroups:
        - !Ref InstanceSecurityGroup

  InstanceSecurityGroup:
    Type: 'AWS::EC2::SecurityGroup'
    Properties:
      GroupDescription: Allow SSH and HTTP access
      SecurityGroupIngress:
        - IpProtocol: tcp
          FromPort: 22
          ToPort: 22
          CidrIp: 0.0.0.0/0
        - IpProtocol: tcp
          FromPort: 80
          ToPort: 80
          CidrIp: 0.0.0.0/0

Outputs:
  InstanceId:
    Description: The ID of the EC2 instance
    Value: !Ref MyInstance

```

## Key Elements in the Template

- **AWSTemplateFormatVersion:** Specifies the version of the template format.
- **Description:** A description of the template.
- **Parameters:** Defines input parameters that can be customized when the stack is created or updated.
- **Resources:** Specifies the AWS resources to be created. In this case:
  - `MyInstance` is an EC2 instance.
  - `InstanceSecurityGroup` is a security group with rules to allow SSH and HTTP access.
- **Outputs:** Defines values that are returned whenever you view your stack's properties. In this case, it returns the instance ID.

## Reading and Writing YAML for CloudFormation

- **Key-Value Pairs:** Simple and intuitive, making YAML easy to read.
- **Nested Objects:** Indentation is used to represent nested objects, making the structure clear.
- **Arrays:** Represented by a dash `-`, which is straightforward to understand.
- **Comments:** Helpful for adding explanations or annotations within the template.

Understanding these basics of YAML will make reading and writing CloudFormation templates more manageable and enhance your ability to leverage AWS infrastructure as code effectively.

# CloudFormation Resources Overview

---

Resources are the core of CloudFormation templates and the only mandatory section. These resources represent various AWS components that are created and configured through the templates. AWS handles the creation, updates, and deletion of these resources based on the template specifications. With over 700 resource types available, understanding how to read the documentation is crucial for effectively using CloudFormation.

## Resource Type Identifiers

Resource types in CloudFormation follow the format: `service-provider::service-name::data-type-name`. This standard helps in identifying and using different AWS resources.

## Finding Resource Documentation

To explore and understand the different types of resources available, you can refer to the [AWS CloudFormation Resource Types Reference](#).

## Example Resources and Documentation

### Example: Amazon EC2

#### 1. EC2 Instance

- **Resource Type:** `AWS::EC2::Instance`
- **Properties:**
  - `AvailabilityZone`: Specifies the Availability Zone.
  - `ImageId`: The AMI ID of the instance.
  - `InstanceType`: The type of instance (e.g., t2.micro).
  - `SecurityGroups`: A list of security group names.
  - **Documentation:** [AWS::EC2::Instance](#)
- **Property Details:**
  - `IamInstanceProfile`: Name of an IAM instance profile (optional, String type, no interruption required if changed).
  - `ImageId`: Requires replacement if changed.

#### 2. Security Groups

- **Resource Type:** `AWS::EC2::SecurityGroup`
- **Properties:**
  - `GroupDescription`: Description of the security group.
  - `SecurityGroupIngress`: Rules for inbound traffic.
  - **Documentation:** [AWS::EC2::SecurityGroup](#)
- **Property Example:**

```
SecurityGroupIngress:
  - IpProtocol: tcp
```

```
FromPort: 22
ToPort: 22
CidrIp: 0.0.0.0/0
- IpProtocol: tcp
FromPort: 80
ToPort: 80
CidrIp: 0.0.0.0/0
```

### Example: Elastic IP

- **Resource Type:** `AWS::EC2::EIP`
- **Properties:**
  - Defines how to declare an elastic IP.
  - **Documentation:** [AWS::EC2::EIP](#)

## Using Documentation for Resource Properties

### 1. Finding Syntax and Examples:

- Each resource documentation provides syntax in JSON and YAML formats.
- Detailed explanations for each property are provided, including whether they are required and the type of value expected.

### 2. Understanding Property Requirements:

- Example for `IamInstanceProfile`: Optional, String type, no interruption required.
- Example for `ImageId`: Requires replacement on change.

## FAQ for Resources

### 1. Creating a Dynamic Number of Resources:

- Use CloudFormation Macros and Transform for dynamic resource creation. This topic is beyond basic usage.

### 2. AWS Service Support:

- Almost all AWS services are supported. For unsupported services, use CloudFormation Custom Resources.

## Conclusion

Understanding how to declare and configure resources in CloudFormation is crucial for automating AWS infrastructure. By leveraging the extensive documentation and following best practices, you can effectively use CloudFormation to manage and deploy AWS resources.



# CloudFormation Parameters Overview

---

Parameters in CloudFormation templates allow users to provide inputs, making templates reusable and adaptable across different environments. They help avoid hardcoding values and provide flexibility for varying configurations.

## When to Use Parameters

1. **Future Configuration Changes:** If a resource configuration is likely to change, use a parameter to avoid re-uploading the template.
2. **Unknown Inputs:** Use parameters when inputs cannot be predetermined.

## Parameter Settings

### 1. Type:

- `String`
- `Number`
- `CommaDelimitedList`
- `List<Number>`
- `AWS-Specific Parameter`
- `SSM Parameter`

2. **Description:** Describes the parameter.

3. **ConstraintDescription:** Provides a description for constraints.

4. **MinLength and MaxLength:** Sets minimum and maximum length for strings.

5. **MinValue and MaxValue:** Sets minimum and maximum values for numbers.

6. **Default:** Specifies a default value.

7. **AllowedValues:** Specifies a list of allowed values.

8. **AllowedPattern:** Uses regex to define allowed patterns.

9. **NoEcho:** Masks sensitive information, such as passwords, from logs.

## Example Parameter: AllowedValues

```
Parameters:
  InstanceType:
    Type: String
    AllowedValues:
      - t2.micro
      - t2.small
      - t2.medium
    Default: t2.micro
```

This parameter allows users to choose from predefined EC2 instance types, ensuring valid input while providing a dropdown selection.

### Example Parameter: NoEcho

```
Parameters:
  DBPassword:
    NoEcho: true
    Type: String
    MinLength: 8
    Description: "The database admin account password"
```

This parameter masks the input to maintain the confidentiality of sensitive information like passwords.

### Using Parameters with !Ref

The `!Ref` function references parameters within the template, enabling their use in resource properties.

#### Example Usage in Template:

```
Parameters:
  SecurityGroupDescription:
    Description: "Description for the security group"
    Type: String

Resources:
  MySecurityGroup:
    Type: AWS::EC2::SecurityGroup
    Properties:
      GroupDescription: !Ref SecurityGroupDescription
```

### Referencing Resources with !Ref

The `!Ref` function can also reference resources, ensuring consistency and correctness within the template.

#### Example:

```
Resources:
  MyInstance:
    Type: AWS::EC2::Instance
    Properties:
      SecurityGroups:
        - !Ref SSHSecurityGroup

  SSHSecurityGroup:
    Type: AWS::EC2::SecurityGroup
```

```
Properties:
  GroupDescription: "Allow SSH access"
```

## Pseudo Parameters

Pseudo parameters are predefined by AWS and can be used directly in templates without declaration.

### Common Pseudo Parameters:

1. **AWS::AccountId:** Returns the AWS account ID.
2. **AWS::Region:** Returns the AWS region in which the stack is deployed.
3. **AWS::StackId:** Returns the stack ID.
4. **AWS::StackName:** Returns the stack name.
5. **AWS::NotificationARNs:** Returns the list of ARNs for the stack's notification topics.
6. **AWS::NoValue:** Specifies that no value is returned.

### Example Usage:

```
Resources:
  MyBucket:
    Type: AWS::S3::Bucket
    Properties:
      BucketName: !Sub
        - "${AWS::AccountId}-${AWS::Region}-my-bucket"
```

This example uses `AWS::AccountId` and `AWS::Region` to dynamically generate a bucket name based on the account ID and region.

## Conclusion

Parameters in CloudFormation templates offer significant flexibility and control, enabling dynamic and reusable infrastructure definitions. By leveraging parameters and pseudo parameters, templates can be customized and adapted for different use cases and environments, ensuring consistency and reducing the potential for errors.

# Mappings in CloudFormation

---

Mappings in CloudFormation templates are used to define fixed variables. They are particularly useful for specifying different values based on specific environments or regions. Here's a detailed explanation:

## Purpose of Mappings

Mappings provide a way to define values that differ based on environment-specific or region-specific variables. They are useful for:

- Differentiating between development (dev) and production (prod) environments.
- Specifying values for different AWS regions.
- Providing architecture-specific settings, such as different AMI IDs for various regions and architectures.

## Structure of Mappings

Mappings are defined within the CloudFormation template with a specific format. Here's an example of how a mapping might look:

```
Mappings:
  RegionMap:
    us-east-1:
      HVM64: "ami-0ff8a91507f77f867"
      HVMG2: "ami-0a584ac55a7631c0c"
    us-west-1:
      HVM64: "ami-0bdbb828fd58c52235"
      HVMG2: "ami-066ee5fd4a9ef77f1"
    eu-west-1:
      HVM64: "ami-047bb4163c506cd98"
      HVMG2: "ami-053b3da8a33b0a4e6"
```

In this example, the `RegionMap` mapping contains different AMI IDs for different regions (`us-east-1`, `us-west-1`, `eu-west-1`) and architectures (`HVM64`, `HVMG2`).

## Accessing Mapping Values

To use the values defined in mappings, the `Fn::FindInMap` intrinsic function is used. This function takes three arguments:

1. **Map Name:** The name of the mapping (e.g., `RegionMap`).
2. **Top-Level Key:** The top-level key in the mapping, which can be a region (e.g., `us-east-1`).
3. **Second-Level Key:** The specific key within the top-level key, such as the architecture type (e.g., `HVM64`).

Here's an example of using `Fn::FindInMap` in an EC2 instance definition:

```
Resources:
  MyEC2Instance:
    Type: "AWS::EC2::Instance"
    Properties:
      ImageId: !FindInMap
        - RegionMap
        - !Ref "AWS::Region"
        - HVM64
```

In this example:

- **RegionMap** is the name of the mapping.
- **!Ref "AWS::Region"** dynamically retrieves the current AWS region using a pseudo parameter.
- **HVM64** specifies the architecture.

If the template is deployed in the **us-east-1** region, **!FindInMap** will retrieve the AMI ID **"ami-0ff8a91507f77f867"** for the **HVM64** architecture.

## When to Use Mappings vs. Parameters

- **Mappings** are best used when:
  - You know all possible values in advance.
  - Values can be deduced from fixed variables like region, environment, or account.
  - You need strict control over the values in your template.
- **Parameters** are best used when:
  - Values are not known ahead of time and depend on user input at runtime.
  - You want to give users flexibility to specify their own values.

## Conclusion

Mappings are a powerful feature in CloudFormation that allow for predefining environment-specific and region-specific values. They provide a way to manage configurations efficiently, ensuring that resources are appropriately configured based on the deployment context. By using mappings, you can avoid hardcoding values and maintain cleaner, more adaptable templates.

# Outputs in CloudFormation

---

The outputs section in a CloudFormation template is optional but very useful. Outputs allow you to declare values that can be imported into other stacks, making it possible to share data and resources across different CloudFormation stacks. Here's a detailed explanation of how outputs work and their benefits.

## Purpose of Outputs

1. **Sharing Data Across Stacks:** Outputs enable you to export values such as VPC IDs or Subnet IDs from one stack and reference them in another. This is particularly useful for linking a network stack with an application stack.
2. **Viewing Values:** You can view the values of outputs in the AWS Management Console or by using the AWS Command Line Interface (CLI). This can be helpful for debugging and management purposes.
3. **Collaboration:** Outputs facilitate collaboration between different teams or experts handling their respective stacks. For example, a network team can create and manage a VPC stack, and an application team can reference the VPC ID in their application stack.

## Example of Outputs

Here's a sample output declaration in a CloudFormation template:

```
Outputs:
  SSHSecurityGroup:
    Description: "The security group ID for SSH access"
    Value: !Ref SSHSecurityGroup
    Export:
      Name: SSHSecurityGroup
```

In this example:

- **Description:** Provides a brief description of what the output represents.
- **Value:** Uses the `!Ref` intrinsic function to reference the security group created within the template.
- **Export:** Specifies a unique name for the output value that can be referenced by other stacks.

## Reusing Outputs

To reuse an exported output in another CloudFormation stack, you use the `Fn::ImportValue` function. Here's an example:

```
Resources:
  MyEC2Instance:
    Type: "AWS::EC2::Instance"
    Properties:
      ImageId: "ami-0ff8a91507f77f867"
```

```
InstanceType: "t2.micro"  
SecurityGroups:  
  - !ImportValue SSHSecurityGroup
```

In this example:

- **!ImportValue SSHSecurityGroup:** Imports the security group ID exported by another stack with the name `SSHSecurityGroup`.

## Important Considerations

1. **Uniqueness:** The name specified in the `Export` block must be unique across all exports in a specific AWS region.
2. **Deletion Dependency:** You cannot delete a stack that has exported values being referenced by other stacks. You must first remove all references to the exported values before deleting the stack.

## Benefits of Outputs

- **Flexibility:** Outputs provide flexibility in organizing and managing AWS resources. Different teams can manage their stacks independently while still being able to reference and use resources created by others.
- **Reusability:** By exporting common resources like VPC IDs or Security Group IDs, you avoid hardcoding values and make your templates more reusable and easier to maintain.
- **Clarity:** Outputs give a clear way to document and share important resource identifiers, making it easier for others to understand and use the stack.

## Conclusion

Outputs are a powerful feature in CloudFormation that enhance the modularity and reusability of your templates. By exporting and importing values, you can create complex architectures where resources are shared and managed efficiently across multiple stacks. This facilitates collaboration and ensures that resources are used consistently throughout your AWS environment.

# Conditions in CloudFormation

---

Conditions in CloudFormation templates allow you to control the creation of resources or outputs based on specific conditions. This feature is useful for managing different environments (such as development, testing, and production) or regional differences. Here's a detailed explanation of how conditions work and how to use them effectively.

## Purpose of Conditions

1. **Environment-Based Configuration:** You can specify resources that should only be created in specific environments. For instance, some resources might only be needed in a production environment and not in a development environment.
2. **Regional Variations:** Conditions can help manage variations based on AWS regions, ensuring that resources are only created where needed.
3. **Parameter-Driven Conditions:** Conditions can be based on parameter values provided during stack creation, allowing for dynamic template behavior.

## Defining Conditions

To define a condition, you use intrinsic functions such as `Fn::And`, `Fn::Equals`, `Fn::If`, `Fn::Not`, and `Fn::Or`. Conditions can reference parameter values and mappings.

### Example Definition

```
Conditions:
  CreateProdResources:
    Fn::Equals: [ !Ref EnvType, "prod" ]
```

In this example, `CreateProdResources` is a condition that evaluates to `true` if the `EnvType` parameter is set to `prod`.

## Using Conditions

Once a condition is defined, it can be applied to resources, outputs, or other parts of the template. If the condition evaluates to `true`, the associated resource or output will be created. If it evaluates to `false`, the resource or output will not be created.

### Example Usage

```
Resources:
  MountPoint:
    Type: "AWS::EC2::VolumeAttachment"
    Condition: CreateProdResources
    Properties:
      InstanceId: !Ref MyEC2Instance
```



```
VolumeId: !Ref MyVolume
Device: "/dev/sdh"
```

In this example:

- The **MountPoint** resource will only be created if the **CreateProdResources** condition is **true**.
- If the condition is **false**, the **MountPoint** resource will not be created.

## Key Functions for Conditions

- **Fn::And**: Logical **AND** operation.
- **Fn::Equals**: Compares two values for equality.
- **Fn::If**: Conditionally includes one of two values.
- **Fn::Not**: Logical **NOT** operation.
- **Fn::Or**: Logical **OR** operation.

## Practical Use Cases

1. **Environment Differentiation**: Create resources that are only needed in certain environments (e.g., additional monitoring resources in production).
2. **Cost Management**: Avoid creating expensive resources in development or testing environments.
3. **Compliance and Security**: Ensure that certain security measures or compliance-related resources are only deployed in appropriate environments.

## Conclusion

Conditions in CloudFormation templates provide a powerful way to manage resource creation dynamically based on various factors such as environment, region, or parameter values. By using conditions, you can create flexible and reusable templates that adapt to different deployment scenarios, helping to streamline your infrastructure management and reduce errors.

Conditions are particularly useful for maintaining a single template that can be used across multiple environments with minimal modifications. Understanding and utilizing conditions effectively will enhance your ability to create robust and adaptable CloudFormation templates.

# Intrinsic Functions in CloudFormation

---

Intrinsic functions in CloudFormation are used to perform specific operations within templates, such as referencing resources, manipulating strings, or getting attribute values. Here's a detailed explanation of the key intrinsic functions that are essential to know, especially for exam purposes.

## Key Intrinsic Functions

### 1. `Ref`

- **Purpose:** Returns the value of a parameter or the physical ID of a resource.
- **Usage:**

```
!Ref MyVPC
```

- Here, `MyVPC` can be either a parameter or another resource in the template.

### 2. `Fn::GetAtt`

- **Purpose:** Retrieves an attribute from a resource.
- **Usage:**

```
!GetAtt EC2Instance.AvailabilityZone
```

- `EC2Instance` is the name of the resource, and `AvailabilityZone` is the attribute.

### 3. `Fn::FindInMap`

- **Purpose:** Retrieves values from a mapping.
- **Usage:**

```
!FindInMap [ RegionMap, !Ref "AWS::Region", "HVM64" ]
```

- Here, `RegionMap` is the mapping name, `!Ref "AWS::Region"` gets the current region, and `HVM64` is the specific key.

### 4. `Fn::ImportValue`

- **Purpose:** Imports values that are exported from other stacks.
- **Usage:**

```
!ImportValue SSHSecurityGroup
```

- This imports the value of `SSHSecurityGroup` exported from another stack.

## 5. `Fn::Join`

- **Purpose:** Joins a list of values into a single string.
- **Usage:**

```
!Join [ ":", [ "a", "b", "c" ] ]
```

- This joins the list `["a", "b", "c"]` into the string `"a:b:c"`.

## 6. `Fn::Sub`

- **Purpose:** Substitutes variables within a string.
- **Usage:**

```
!Sub "arn:aws:s3:::${BucketName}/*"
```

- Here, `${BucketName}` will be replaced with the value of the `BucketName` parameter or resource.

## 7. `Fn::Base64`

- **Purpose:** Converts a string into its Base64 representation.
- **Usage:**

```
!Base64 "Hello, World!"
```

- Converts `"Hello, World!"` to its Base64 equivalent.

## 8. `Fn::Cidr`

- **Purpose:** Returns an array of CIDR blocks.
- **Usage:**

```
!Cidr [ !Ref "VpcCidrBlock", 6, 5 ]
```

- Generates CIDR blocks based on the VPC CIDR block.

## 9. `Fn::GetAZs`

- **Purpose:** Returns an array of Availability Zones for a given region.
- **Usage:**

```
!GetAZs "us-east-1"
```

- Returns the list of AZs in the `us-east-1` region.

## 10. `Fn::Select`

- **Purpose:** Selects an element from a list.
- **Usage:**

```
!Select [ 1, [ "a", "b", "c" ] ]
```

- Selects the element at index 1, which is `"b"`.

## 11. `Fn::Split`

- **Purpose:** Splits a string into a list of strings.
- **Usage:**

```
!Split [ ",", "a,b,c" ]
```

- Splits the string `"a,b,c"` into the list `["a", "b", "c"]`.

## 12. `Fn::Transform`

- **Purpose:** Applies a macro to process the template.
- **Usage:**

```
Fn::Transform:
  Name: "AWS::Include"
  Parameters:
    Location: "s3://my-bucket/my-template.yml"
```

- Applies the `AWS::Include` transform to include an external template.

## Condition Functions

These functions are used within conditions to control resource creation based on logical operations.

### 1. `Fn::And`

- **Purpose:** Logical **AND** operation.
- **Usage:**

```
!And [ condition1, condition2 ]
```

## 2. **Fn::Equals**

- **Purpose:** Compares two values.
- **Usage:**

```
!Equals [ value1, value2 ]
```

## 3. **Fn::If**

- **Purpose:** Conditional value selection.
- **Usage:**

```
!If [ condition, valueIfTrue, valueIfFalse ]
```

## 4. **Fn::Not**

- **Purpose:** Logical **NOT** operation.
- **Usage:**

```
!Not [ condition ]
```

## 5. **Fn::Or**

- **Purpose:** Logical **OR** operation.
- **Usage:**

```
!Or [ condition1, condition2 ]
```

## Practical Examples

### Ref and GetAtt Example

```

Resources:
  MyEC2Instance:
    Type: "AWS::EC2::Instance"
    Properties:
      ImageId: "ami-0ff8a91507f77f867"
      InstanceType: "t2.micro"

  MyVolume:
    Type: "AWS::EC2::Volume"
    Properties:
      AvailabilityZone: !GetAtt MyEC2Instance.AvailabilityZone
      Size: 100

```

- `MyVolume` uses `!GetAtt MyEC2Instance.AvailabilityZone` to set its `AvailabilityZone` property based on the instance's zone.

## FindInMap Example

```

Mappings:
  RegionMap:
    us-east-1:
      HVM64: "ami-0ff8a91507f77f867"
    us-west-1:
      HVM64: "ami-0bdb828fd58c52235"

Resources:
  MyEC2Instance:
    Type: "AWS::EC2::Instance"
    Properties:
      ImageId: !FindInMap [ RegionMap, !Ref "AWS::Region", "HVM64" ]

```

- The `ImageId` is dynamically selected based on the current region using `!FindInMap`.

## Conclusion

Understanding and effectively using intrinsic functions in CloudFormation is crucial for creating dynamic, reusable, and efficient templates. These functions enable complex configurations, conditional resource creation, and inter-stack references, greatly enhancing the flexibility and power of your CloudFormation templates.

# CloudFormation Rollbacks

---

CloudFormation rollbacks are crucial for maintaining the integrity and consistency of your stacks during creation and update processes. Here's a detailed explanation of how rollbacks work, options available, and how to manage failures effectively.

## Stack Creation Failures

When you create a stack and it fails, there are two primary options for handling the rollback:

### 1. Default Rollback:

- If the stack creation fails, CloudFormation rolls back and deletes all resources that were created as part of that stack.
- This ensures that no partial resources are left behind, but it also means you cannot inspect the resources to troubleshoot the failure.

### 2. Disable Rollback:

- You can disable the rollback to keep the resources created up to the point of failure.
- This allows you to inspect and troubleshoot what went wrong.
- This option can be set during stack creation under Stack failure options by selecting "Preserve successfully provisioned resources."

## Stack Update Failures

When updating a stack, CloudFormation handles failures by rolling back to the last known stable state:

### 1. Automatic Rollback:

- By default, if a stack update fails, CloudFormation will roll back to the previous known working state.
- This means any new resources or changes that were part of the update will be deleted.

### 2. ContinueUpdateRollback:

- If a rollback itself fails, you may have to intervene manually to fix the resources that were causing issues.
- After manual intervention, you can issue a `ContinueUpdateRollback` command to resume the rollback process.
- This can be done through the console, API, or CLI using the `ContinueUpdateRollback` API call.

## Practical Example

### 1. Creating a Stack with a Known Failure:

- Use a template with an invalid EC2 instance image ID to trigger a failure.

- Example Template (`trigger-failure.yaml`):

```
Resources:
  MyInstance:
    Type: "AWS::EC2::Instance"
    Properties:
      ImageId: "ami-invalid"
      InstanceType: "t2.micro"
```

- When creating the stack, select "Preserve successfully provisioned resources" under Stack failure options.
- This will keep any resources created before the failure, allowing you to inspect them.

## 2. Updating a Stack and Triggering a Failure:

- Create a stack using a valid template (`just-ec2.yaml`), then update it using the template with the invalid EC2 instance image ID.
- During the update, select the default rollback option to see how CloudFormation deletes all resources created as part of the update.

## 3. Handling Rollback Failures:

- If a rollback fails, identify and fix the manually changed resources.
- Use the `ContinueUpdateRollback` command to resume the rollback process.

## Summary

Understanding and managing CloudFormation rollbacks is essential for maintaining your AWS infrastructure's stability and consistency. The key points include:

- **Default Rollback:** Automatically deletes all resources if stack creation fails.
- **Disable Rollback:** Keeps resources created before failure for troubleshooting.
- **Automatic Update Rollback:** Rolls back to the previous stable state on update failure.
- **ContinueUpdateRollback:** Allows resuming a failed rollback after manual intervention.

These mechanisms help ensure that your stacks remain consistent and that failures can be managed effectively.



# CloudFormation and Security: Using Service Roles

---

In AWS CloudFormation, service roles are IAM roles that allow CloudFormation to perform operations on your behalf. This ensures that you can manage stack resources securely and with the principle of least privilege.

## What are Service Roles?

Service roles are IAM roles specifically created for AWS CloudFormation. They allow CloudFormation to create, update, and delete stack resources without needing to grant the same permissions directly to users.

## Use Case

- **Least Privilege Principle:** You can restrict user permissions while still allowing them to manage CloudFormation stacks. Users only need permissions to interact with CloudFormation and to pass the service role.

## Steps to Create and Use Service Roles

### 1. Create the IAM Role for CloudFormation:

- Go to the IAM console.
- Navigate to the roles section and create a new role.
- Select "AWS service" and then "CloudFormation."
- Attach the necessary permissions to this role, for example, [AmazonS3FullAccess](#).
- Name the role (e.g., [DemoRoleForCFNWithS3Capabilities](#)).

### 2. Assign the Service Role to a CloudFormation Stack:

- When creating or updating a stack in CloudFormation, specify the IAM role in the "Permissions" section.
- Select the service role created in the previous step ([DemoRoleForCFNWithS3Capabilities](#)).
- CloudFormation will use this role to perform operations instead of using the user's permissions.

## Example Scenario

### 1. Create IAM Role:

- Create a role for the CloudFormation service.
- Attach [AmazonS3FullAccess](#) to the role.
- Name it [DemoRoleForCFNWithS3Capabilities](#).

### 2. Use the Role in a CloudFormation Stack:

- Go to the CloudFormation console.
- Create a new stack and select an existing template.

- In the "Permissions" section, specify the IAM role (`DemoRoleForCFNWithS3Capabilities`).
- CloudFormation will attempt to create resources using the permissions granted by this role.

If the template attempts to create resources beyond the permissions of the role (e.g., EC2 instances), the stack operation will fail due to insufficient permissions.

### Key Points

- **IAM PassRole Permission:** Users must have the `iam:PassRole` permission to pass a service role to CloudFormation.
- **Service Role Scope:** The role defines the scope of operations CloudFormation can perform, ensuring tighter security and adherence to the principle of least privilege.

Using service roles enhances security by ensuring that users do not need extensive permissions to manage stack resources directly. Instead, CloudFormation performs these operations within the boundaries defined by the service role.

# CloudFormation Capabilities

---

When using AWS CloudFormation to create or update stacks, especially those involving IAM resources, it is crucial to specify certain capabilities to acknowledge the creation and management of these resources. Here's a detailed explanation of the key capabilities you need to understand:

## **CAPABILITY\_NAMED\_IAM and CAPABILITY\_IAM**

- **CAPABILITY\_NAMED\_IAM:** This capability is required when your CloudFormation template is creating or updating IAM resources with custom names. For example, if your template includes an IAM user, role, group, or policy with specified names, you need to specify this capability. It explicitly acknowledges that named IAM resources are being created.
- **CAPABILITY\_IAM:** This capability is necessary when your template includes the creation or update of IAM resources but does not specify custom names for these resources. This acknowledgment is required for any IAM-related operations within the template.

The purpose of these capabilities is to ensure that you are fully aware and explicitly acknowledge that IAM resources, which can significantly impact security and access management, are being managed by CloudFormation.

## **CAPABILITY\_AUTO\_EXPAND**

- **CAPABILITY\_AUTO\_EXPAND:** This capability is used when your CloudFormation template includes macros or nested stacks (stacks within stacks). It acknowledges that the template may dynamically transform before deployment. This is important for complex templates that perform dynamic operations and need to expand or transform other templates.

## **Handling InsufficientCapabilitiesException**

- **InsufficientCapabilitiesException:** If you encounter this exception while launching a template, it indicates that the required capabilities were not acknowledged. To resolve this, you must redo the template launch and explicitly specify the required capabilities. This can be done by adding an extra argument in your API call or by checking the appropriate box in the AWS Management Console.

## Example Scenario

Here's an example scenario demonstrating the use of these capabilities:

1. **Create IAM Role Template:** A CloudFormation template (`capabilities.yaml`) is designed to create an IAM role with a custom name (`MyCustomRoleName`) and a managed policy (`AmazonEC2FullAccess`).
2. **Launching the Stack:**
  - Upload the template file (`capabilities.yaml`).
  - When creating the stack, acknowledge the required capabilities.

### 3. Acknowledging Capabilities:

- During the stack creation process in the AWS Management Console, you will be prompted to acknowledge the creation of IAM resources with custom names.
- Check the box indicating you understand and acknowledge this requirement.

Here is an example of how this process works in practice:

```
AWSTemplateFormatVersion: '2010-09-09'
Resources:
  MyCustomRole:
    Type: "AWS::IAM::Role"
    Properties:
      RoleName: "MyCustomRoleName"
      AssumeRolePolicyDocument:
        Version: "2012-10-17"
        Statement:
          - Effect: "Allow"
            Principal:
              Service:
                - "ec2.amazonaws.com"
            Action:
              - "sts:AssumeRole"
      ManagedPolicyArns:
        - "arn:aws:iam::aws:policy/AmazonEC2FullAccess"
```

In the AWS Management Console:

- Go through the stack creation steps.
- Under the permissions section, you will see a prompt to acknowledge the IAM capabilities.
- Check the acknowledgment box for `CAPABILITY_NAMED_IAM`.

If the acknowledgment is not provided, the stack creation will fail with an error indicating the lack of necessary capabilities.

### Summary

Understanding and properly using CloudFormation capabilities ensures that you explicitly acknowledge and manage the creation and updating of IAM resources, which is critical for maintaining secure and controlled access to AWS resources.

# CloudFormation DeletionPolicy

---

The **DeletionPolicy** attribute in AWS CloudFormation allows you to control the fate of resources when they are removed from a template or when the stack is deleted. This feature is crucial for preserving and backing up important resources.

## Default DeletionPolicy

- **Delete:** By default, when a CloudFormation stack is deleted, all resources within it are also deleted. This behavior can be explicitly specified using **DeletionPolicy: Delete**.

Example:

```
Resources:
  MyEC2Instance:
    Type: "AWS::EC2::Instance"
    DeletionPolicy: Delete
```

In this case, the EC2 instance will be deleted when the stack is deleted.

## Handling S3 Buckets

For S3 buckets, there is a special consideration. If an S3 bucket is not empty, specifying **DeletionPolicy: Delete** will fail unless the bucket is manually emptied or a custom resource is implemented to clear the bucket contents before deletion.

## Retain DeletionPolicy

- **Retain:** This policy is used to preserve resources even after the CloudFormation stack is deleted. It is useful when you want to keep the data or configuration in resources like databases or security groups.

Example:

```
Resources:
  MyDynamoDBTable:
    Type: "AWS::DynamoDB::Table"
    DeletionPolicy: Retain
```

In this case, the DynamoDB table will be retained even if the stack is deleted.

## Snapshot DeletionPolicy

- **Snapshot:** This policy creates a final snapshot of the resource before deletion. It is particularly useful for resources that support snapshots, such as EBS volumes, RDS instances, and other

database services.

Example:

```
Resources:
  MyEBSVolume:
    Type: "AWS::EC2::Volume"
    DeletionPolicy: Snapshot
```

Here, a snapshot of the EBS volume will be taken before it is deleted.

### Example Scenario

Consider a CloudFormation template with a security group and an EBS volume, demonstrating different **DeletionPolicy** settings:

```
Resources:
  MySecurityGroup:
    Type: "AWS::EC2::SecurityGroup"
    DeletionPolicy: Retain

  MyEBSVolume:
    Type: "AWS::EC2::Volume"
    DeletionPolicy: Snapshot
```

#### 1. Stack Creation:

- Create a stack using the above template.
- The stack will include an EBS volume and an EC2 security group.

#### 2. Stack Deletion:

- When the stack is deleted, the EBS volume will be deleted, but a snapshot will be created before its deletion.
- The security group will not be deleted due to the **Retain** policy.

#### 3. Verification:

- After the stack deletion, check the CloudFormation events to verify that the security group deletion was skipped.
- Confirm the creation of the EBS volume snapshot.

### Summary

The **DeletionPolicy** attribute is a powerful feature in CloudFormation for managing the lifecycle of resources. It allows you to:

- Ensure the default deletion of resources.

- Retain resources for future use.
- Create snapshots for backup before deletion.

By using these policies effectively, you can maintain data integrity and manage resource cleanup in a controlled manner.

# CloudFormation Stack Policies

---

CloudFormation Stack Policies are JSON documents that define what update actions are allowed on specific resources during stack updates. They are used to protect critical resources from unintentional updates.

## Default Behavior

- **Unrestricted Updates:** By default, any action is allowed on all resources in a CloudFormation stack. This means you can modify any part of your stack during an update.

## Using Stack Policies

- **Protection:** Stack policies help protect critical parts of your stack against updates. This is especially useful for resources like production databases where unintended changes could cause significant issues.

## Example of a Stack Policy

Here is a sample JSON document for a stack policy:

```
{
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "Update:*",
      "Resource": "*"
    },
    {
      "Effect": "Deny",
      "Action": "Update:*",
      "Resource": "arn:aws:cloudformation:region:account-id:stack/stack-
name/ProductionDatabase"
    }
  ]
}
```

- **Allow Statement:** The first statement allows updates (**Update:\***) on all resources (**Resource: \***) in the stack.
- **Deny Statement:** The second statement denies updates (**Update:\***) specifically on the **ProductionDatabase** resource.

## Goals of Stack Policies

- **Protection:** Stack policies are primarily used to protect resources from unintentional updates.



- **Explicit Allowances:** When setting a stack policy, all resources are protected by default, and you need to explicitly allow updates for specific resources you want to change.

## Implementing a Stack Policy

1. **Create or update a stack:** You can set a stack policy when creating or updating a stack.
2. **Specify the policy:** Use the AWS Management Console, AWS CLI, or AWS SDKs to attach the policy to your stack.

Example using AWS CLI:

```
aws cloudformation set-stack-policy --stack-name my-stack --stack-policy-body  
file://stack-policy.json
```

This command applies the policy defined in `stack-policy.json` to `my-stack`.

## Summary

CloudFormation Stack Policies provide a mechanism to control and restrict updates to specific resources within a stack. They are particularly useful for protecting critical components like production databases from unintended modifications. By default, all resources are protected when a stack policy is in place, and explicit allowances must be specified for resources that can be updated.

# CloudFormation Termination Protection

---

Termination Protection is a feature in AWS CloudFormation that prevents accidental deletion of your CloudFormation stacks. Enabling this feature ensures that a stack cannot be deleted unless the protection is explicitly disabled.

## Steps to Enable Termination Protection

### 1. Create a Stack:

- Go to the AWS CloudFormation console.
- Create a new stack by uploading a template file (e.g., `just-ec2.yaml`).
- Provide a stack name (e.g., `demo`), and proceed through the setup steps.

### 2. Enable Termination Protection:

- Once the stack is created, navigate to the stack's settings.
- Locate the option for Termination Protection. By default, this is deactivated.
- Activate Termination Protection.

### 3. Verify Protection:

- With Termination Protection enabled, any attempt to delete the stack will be blocked.
- The console will display a message indicating that Termination Protection is enabled and must be disabled before the stack can be deleted.

### 4. Disable Termination Protection (if necessary):

- If you have the necessary permissions, you can deactivate Termination Protection.
- After deactivating it, you can proceed to delete the stack.

## Example Scenario

- **Creation:** You create a stack using a template (`just-ec2.yaml`), name it `demo`, and enable Termination Protection.
- **Protection:** Any attempts to delete the `demo` stack will fail with a message indicating that Termination Protection is enabled.
- **Disabling:** If you decide to delete the stack, you must first disable Termination Protection. After disabling, you can delete the stack.

## Importance of Termination Protection

- **Prevent Accidental Deletion:** Termination Protection is a safeguard to prevent accidental deletions, which can be crucial for production environments.
- **Controlled Deletion:** Ensures that only users with the necessary permissions can disable the protection and delete the stack, adding a layer of security.

## Summary

Termination Protection is a simple yet effective feature in CloudFormation that prevents accidental deletion of stacks. It is especially useful for protecting critical resources and ensuring that deletions are intentional and controlled. To delete a protected stack, Termination Protection must first be disabled by a user with the appropriate permissions.

# CloudFormation Custom Resources

---

CloudFormation supports numerous AWS resources, but there are cases where you need to manage resources that are not yet supported or require custom provisioning logic. Custom resources in CloudFormation allow you to handle these scenarios effectively.

## Use Cases for Custom Resources

1. **Unsupported AWS Resources:** Manage AWS resources that are not natively supported by CloudFormation.
2. **External Resources:** Provision and manage resources outside of AWS, such as on-premises or third-party resources.
3. **Custom Logic:** Implement custom scripts during the create, update, and delete phases of your CloudFormation stack using AWS Lambda functions.

## Example Use Case: Emptying an S3 Bucket Before Deletion

A common scenario where custom resources are used is emptying an S3 bucket before deleting it, as CloudFormation cannot delete a non-empty S3 bucket. Here's how you can achieve this using a Lambda function.

## Defining a Custom Resource

A custom resource in CloudFormation is defined in the template with the type `Custom::MyCustomResourceTypeName`. It is typically backed by an AWS Lambda function or an SNS topic, though Lambda is more common.

Here's the basic structure for defining a custom resource:

```
Resources:
  MyCustomResource:
    Type: Custom::MyLambdaResource
    Properties:
      ServiceToken: arn:aws:lambda:region:account-id:function:MyLambdaFunction
      # Additional properties to pass input data to the Lambda function
```

## Key Components

1. **ServiceToken:** The ARN of the Lambda function (or SNS topic) that implements the logic for the custom resource.
2. **Properties:** Any additional parameters needed by the Lambda function to perform its task.

## Example: Custom Resource for Emptying an S3 Bucket

The following example demonstrates a custom resource that triggers a Lambda function to empty an S3 bucket before the stack deletion:

1. **Lambda Function:** This function will contain the logic to list and delete objects from the specified S3 bucket.

## 2. CloudFormation Template:

```
Resources:
  MyS3Bucket:
    Type: AWS::S3::Bucket
    Properties:
      BucketName: my-bucket

  EmptyS3BucketFunction:
    Type: AWS::Lambda::Function
    Properties:
      Handler: index.handler
      Role: arn:aws:iam::account-id:role/lambda-execution-role
      Code:
        ZipFile: |
          import boto3
          import cfnresponse

          def handler(event, context):
              s3 = boto3.client('s3')
              bucket = event['ResourceProperties']['BucketName']

              if event['RequestType'] == 'Delete':
                  objects = s3.list_objects_v2(Bucket=bucket)
                  if 'Contents' in objects:
                      keys = [{'Key': obj['Key']} for obj in
objects['Contents']]
                      s3.delete_objects(Bucket=bucket, Delete={'Objects':
keys})

                  cfnresponse.send(event, context, cfnresponse.SUCCESS, {})

  CustomResourceToEmptyBucket:
    Type: Custom::EmptyS3Bucket
    Properties:
      ServiceToken: !GetAtt EmptyS3BucketFunction.Arn
      BucketName: !Ref MyS3Bucket
```

## Workflow

1. **Creation:** The CloudFormation stack creates the S3 bucket and the Lambda function.
2. **Deletion:** When the stack is deleted, the custom resource triggers the Lambda function, which:
  - Lists objects in the S3 bucket.
  - Deletes the objects from the bucket.
  - Signals CloudFormation to proceed with the deletion of the now-empty S3 bucket.

## Summary

Custom resources in CloudFormation provide flexibility to manage resources and execute custom logic that is not natively supported. By leveraging AWS Lambda, you can handle complex operations such as emptying an S3 bucket before deletion, ensuring smooth and automated resource management within your CloudFormation stacks.

# CloudFormation StackSets

---

CloudFormation StackSets provide a way to manage stacks across multiple AWS accounts and regions with a single operation or template. This is particularly useful for deploying infrastructure consistently across an organization or multiple environments.

## Key Concepts

1. **Administrative Account:** The StackSet is created and managed from an administrative AWS account. This account must have the necessary permissions to manage resources across target accounts and regions.
2. **Template:** A CloudFormation template serves as the blueprint for the resources you want to deploy across multiple accounts and regions.
3. **Target Accounts and Regions:** These are the AWS accounts and regions where you want to deploy your stacks. You can specify one or more AWS accounts and one or more regions per StackSet.
4. **Update Behavior:** When you update a StackSet, CloudFormation ensures that all stack instances across all specified target accounts and regions are updated simultaneously. This ensures consistency and avoids the need to manage each stack individually.
5. **AWS Organizations Integration:** A common use case for StackSets is deploying stacks across all accounts within an AWS Organization. AWS Organizations is a service that helps you centrally manage and govern multiple AWS accounts.

## Creating and Using StackSets

To create a StackSet, you define a CloudFormation template that includes parameters for specifying target accounts and regions. Here's a simplified example of how you might define a StackSet:

```
Resources:
  MyStackSet:
    Type: AWS::CloudFormation::StackSet
    Properties:
      TemplateURL: https://s3.amazonaws.com/my-bucket/my-template.yaml
      PermissionModel: SELF_MANAGED
      StackInstances:
        - DeploymentTargets:
            Accounts:
              - account-id-1
              - account-id-2
            Regions:
              - us-east-1
              - us-west-2
        # More StackInstances for additional accounts and regions can be added
```

## Update and Management

- **Update:** When you update the StackSet with a new template or configuration, CloudFormation applies the changes to all specified stack instances concurrently.
- **Permissions:** Only an administrator or a user with delegated permissions within AWS Organizations can create and manage StackSets. This helps maintain security and control over deployments.

## Benefits of StackSets

- **Consistency:** Ensure consistent deployments of infrastructure across multiple accounts and regions.
- **Centralized Management:** Manage all deployments from a single administrative account.
- **Scale:** Easily scale deployments across hundreds or even thousands of AWS accounts.

## Considerations

- **Permissions:** Ensure that the administrative account has the necessary IAM permissions to manage resources in target accounts.
- **Rollback:** StackSets support automatic rollback in case of deployment failures, ensuring that your infrastructure remains consistent and operational.

## Conclusion

CloudFormation StackSets simplify the management of infrastructure deployments at scale across AWS accounts and regions. By leveraging StackSets, organizations can achieve consistent and efficient infrastructure provisioning while maintaining control and security over their AWS environments.